



Module 7: Hands-On: Creating a Deployment

Service in Kubernetes

Below is the command for Service yaml file:

apiVersion: v1

kind: Service

metadata: name:

my-nginx

spec:

type: NodePort

ports:

- targetPort: 80

port: 80

nodePort: 30008

selector:

app: nginx

```
GNU nano 4.8 nginx_service.yml
apiVersion: v1
kind: Service
metadata:
  name: my-nginx

spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: nginx
```

Note: the key value pair in the selector is taken from the deployment file.

```
GNU nano 4.8 deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	32m
my-nginx	NodePort	10.103.247.61	<none>	80:30008/TCP	15m

To create a service:

```
kubectl create -f nginx_service.yml
```

← → ↻ ⚠ Not secure | 3.142.184.14:30008

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

IntelliPaat