



## **Module 7: Hands-On: Kubernetes Dashboard**

Run the following commands for installing kubeadm as root (both master and worker)

```
sudo su
```

```
apt-get update
```

```
apt-get install docker.io
```

```
apt-get update && apt-get install -y apt-transport-https curl
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
cat <<EOF > /etc/apt/sources.list.d/kubernetes.list
```

```
deb https://apt.kubernetes.io/ kubernetes-xenial main
```

```
EOF
```

```
apt-get update
```

```
apt install kubeadm=1.21.0-00 kubectl=1.21.0-00 kubelet=1.21.0-00 -y
```

## Creating cluster:

Initializing kubeadm on master using:

```
kubeadm init --pod-network-cidr=192.168.0.0/16
```

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.23.197:6443 --token om12ul.nhq0adgipvqii3yg \
--discovery-token-ca-cert-hash sha256:1665c8651063e91650362aa60aa89724d63c68397f34387b44e715d6
90ca6bca
root@ip-172-31-23-197:/home/ubuntu#
```

Copy the Join token and paste in the worker machine.

```
This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

We need to run the below commands:

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

To list all nodes :

```
kubectl get nodes
```

```
root@ip-172-31-23-197:/home/ubuntu# kubectl get nodes
NAME                STATUS    ROLES                  AGE      VERSION
ip-172-31-23-197    NotReady control-plane,master   4m44s    v1.21.0
ip-172-31-37-241    NotReady <none>                4m17s    v1.21.0
```

It shows the nodes but the status is not ready because we have not installed the network plugin.

To install network plugin, run the below commands:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
```

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.49.0/deploy/static/provider/baremetal/deploy.yaml
```

Check the nodes for its state after installing network plugins.

```
kubectl get nodes
```

```
root@ip-172-31-23-197:/home/ubuntu# kubectl get nodes
NAME                STATUS    ROLES                  AGE      VERSION
ip-172-31-23-197    Ready     control-plane,master   6m54s    v1.21.0
ip-172-31-37-241    Ready     <none>                 6m27s    v1.21.0
```

Thus, we have successfully installed Kubernetes. Run the below command to create a dashboard:


```
kubectl apply -f
```

```
https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0/aio/deploy/recommended.yaml
```

Then edit the service:

```
kubectl edit service kubernetes-dashboard -n kubernetes-dashboard
```

```
Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-app":
s-dashboard"},"spec":{"ports":[{"port":443,"targetPort":8443}],"selector":{"k8s-app":"kuber
creationTimestamp: "2022-01-13T07:16:45Z"
labels:
  k8s-app: kubernetes-dashboard
name: kubernetes-dashboard
namespace: kubernetes-dashboard
resourceVersion: "1539"
uid: 01f6f13a-d1bf-45b0-afc4-ca9713168cac
spec:
  clusterIP: 10.105.231.218
  clusterIPs:
  - 10.105.231.218
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - port: 443
    protocol: TCP
    targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: None
  type: ClusterIP
status:
```



Note: We need to change the type from ClusterIP to NodePort. The below given image contains the modified service file.

```
apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-app":"kubernetes-dashboard"},"spec":{"ports":[{"port":443,"targetPort":8443}],"selector":{"k8s-app":"kubernetes-dashboard"},"status":{"loadBalancer":{}}}}
creationTimestamp: "2022-01-13T07:16:45Z"
labels:
  k8s-app: kubernetes-dashboard
name: kubernetes-dashboard
namespace: kubernetes-dashboard
resourceVersion: "1997"
uid: 01f6f13a-d1bf-45b0-afc4-ca9713168cac
spec:
  clusterIP: 10.105.231.218
  clusterIPs:
  - 10.105.231.218
  externalTrafficPolicy: Cluster
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - nodePort: 31707
    port: 443
    protocol: TCP
    targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: None
  type: NodePort
status:
  loadBalancer: {}
```

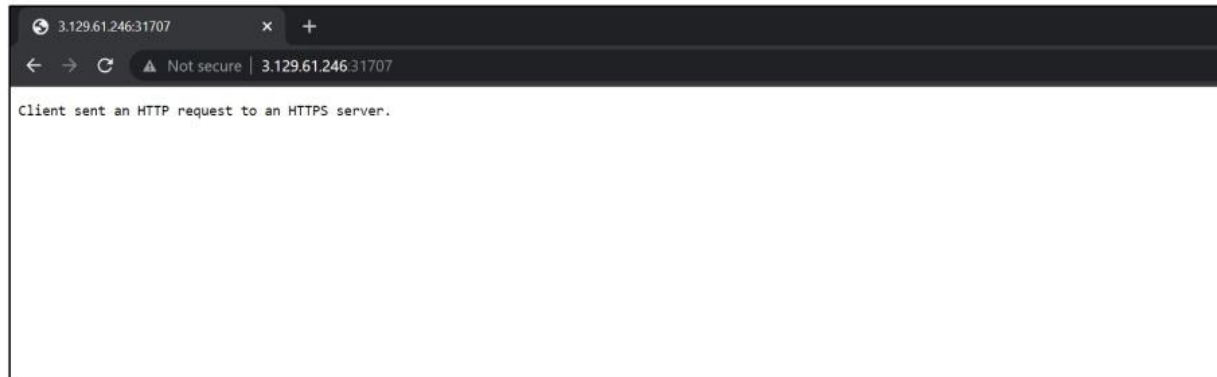
Note: The editor used is vim. Therefore to save and exit we need to press ESC and then: wq

```
root@ip-172-31-23-197:/home/ubuntu# kubectl edit service kubernetes-dashboard -n kubernetes-dashboard
service/kubernetes-dashboard edited
```

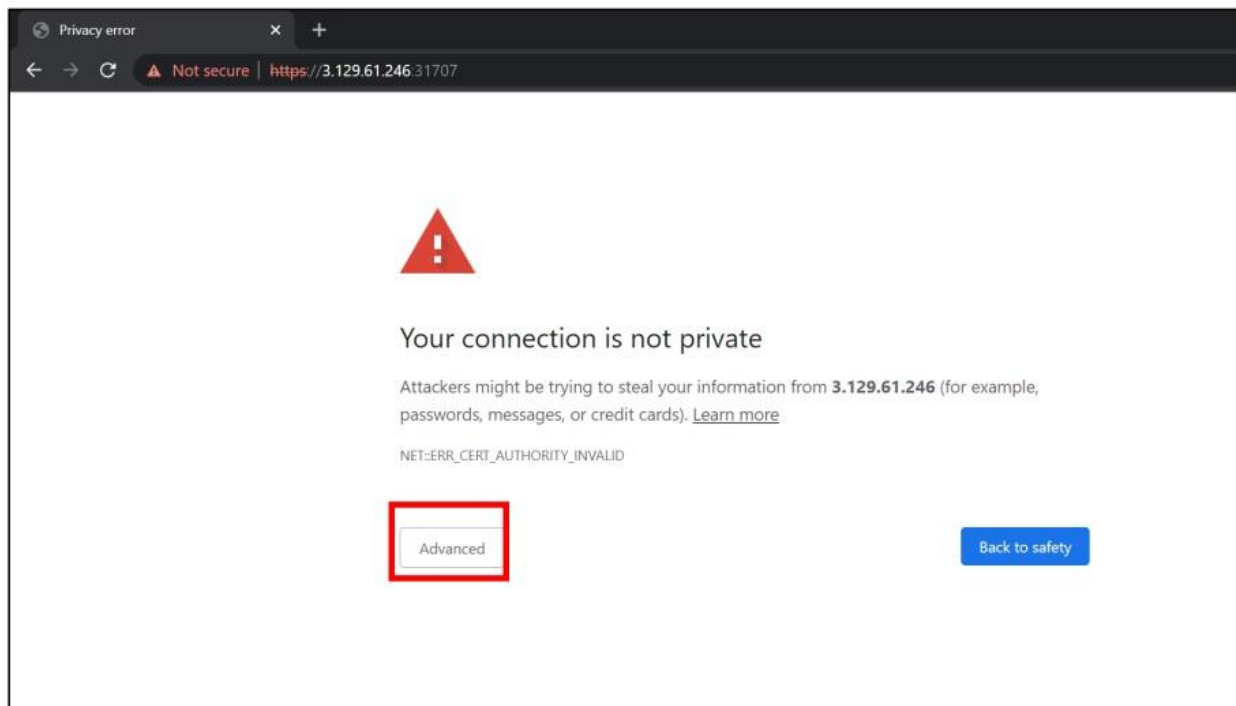
***kubectl get svc -n kubernetes-dashboard***

```
root@ip-172-31-23-197:/home/ubuntu# kubectl get svc -n kubernetes-dashboard
NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
dashboard-metrics-scraper          ClusterIP    10.108.75.42     <none>            8000/TCP         16m
kubernetes-dashboard               NodePort     10.105.231.218   <none>            443:31707/TCP    16m
root@ip-172-31-23-197:/home/ubuntu#
```

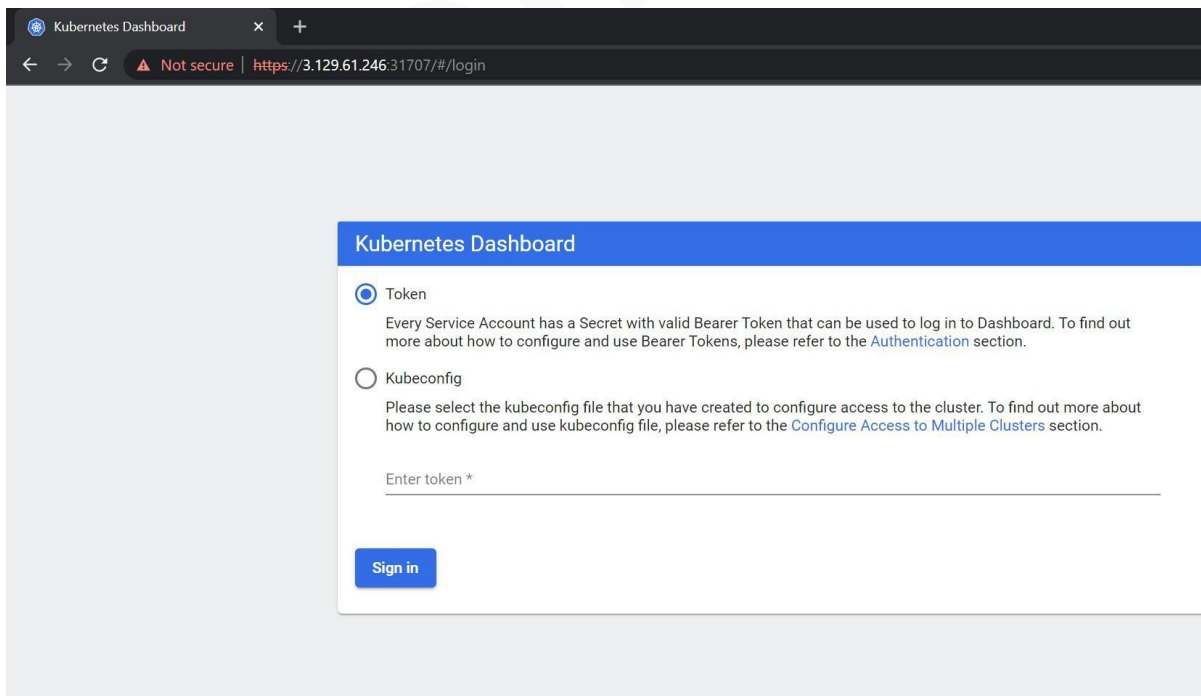
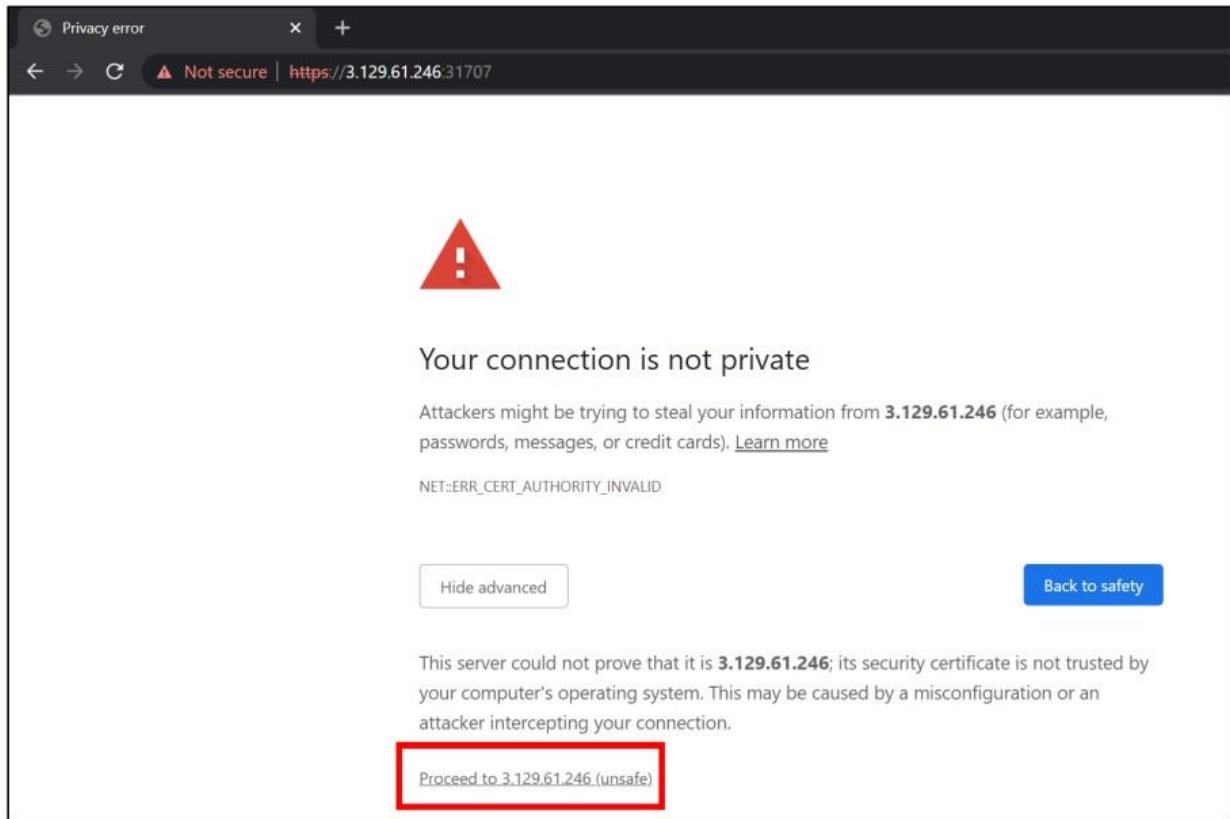
Now go to a browser and paste the ip along with the port  
<https://<ip-of-master>:31707>



We need to click on Advance option to access the webpage



Click on the Proceed option to open the dashboard





We need to create a service account:

To create service account run the below command:

```
kubectl create serviceaccount cluster-admin-dashboard-sa
```

To bind clusterAdmin role to the service account use the below command:

```
kubectl create clusterrolebinding cluster-admin-dashboard-sa \
--clusterrole=cluster-admin \
--serviceaccount=default:cluster-admin-dashboard-sa
```

To parse the token run the below command:

```
TOKEN=$(kubectl describe secret $(kubectl -n kube-system get secret | awk '/^cluster-admin-boardsa-token-/ {print $1}') | awk '$1=="token:" {print $2}')
```

Then we need to run the below command:

***echo \$TOKEN***

[illegible]

Copy the token and paste in the Kubernetes dashboard.



## Kubernetes Dashboard

☒ Token

Every Service Account has a Secret with valid Bearer Token that can be used to log in to Dashboard. To find out more about how to configure and use Bearer Tokens, please refer to the [Authentication](#) section.

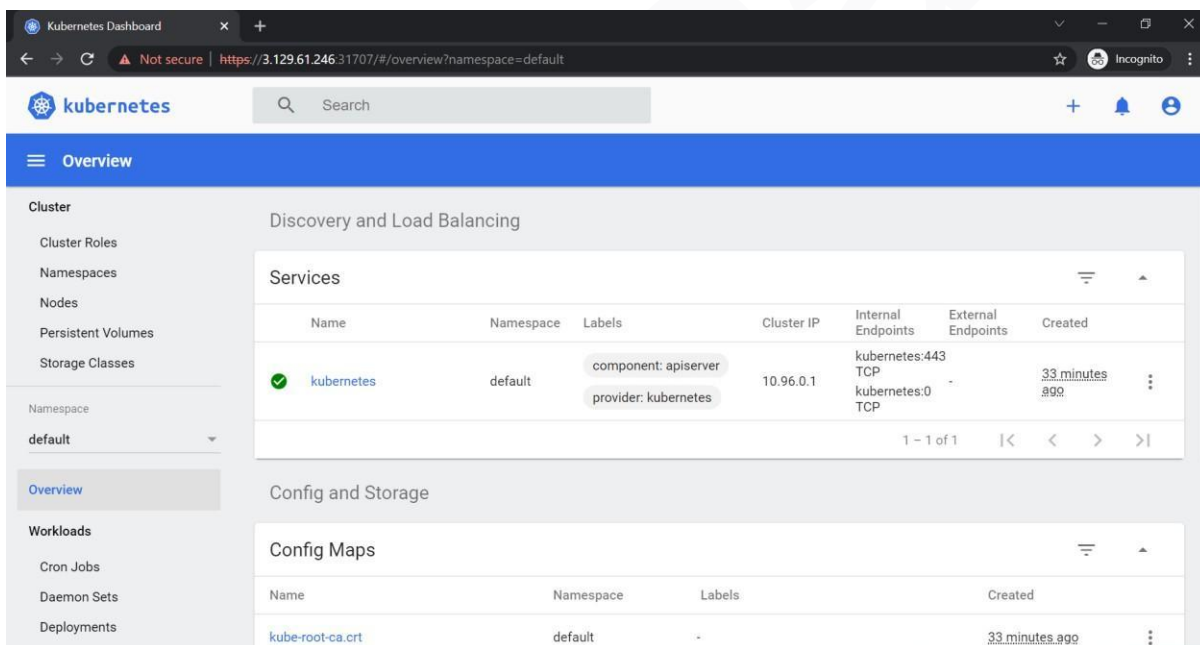
☐ Kubeconfig

Please select the kubeconfig file that you have created to configure access to the cluster. To find out more about how to configure and use kubeconfig file, please refer to the [Configure Access to Multiple Clusters](#) section.

Enter token \*

.....

[Sign in](#)



The screenshot shows the Kubernetes Dashboard in a web browser. The URL is <https://3.129.61.246:31707/#/overview?namespace=default>. The dashboard has a sidebar with navigation links: Overview, Cluster, Cluster Roles, Namespaces, Nodes, Persistent Volumes, Storage Classes, Namespace (set to default), Overview (selected), Workloads, Cron Jobs, Daemon Sets, and Deployments. The main content area is divided into two sections: 'Discovery and Load Balancing' and 'Config and Storage'. The 'Discovery and Load Balancing' section contains a 'Services' table with one entry: 'kubernetes' in the 'default' namespace, with labels 'component: apiserver' and 'provider: kubernetes', and a cluster IP of '10.96.0.1'. The 'Config and Storage' section contains a 'Config Maps' table with one entry: 'kube-root-ca.crt' in the 'default' namespace, created '33 minutes ago'.

Name	Namespace	Labels	Cluster IP	Internal Endpoints	External Endpoints	Created
kubernetes	default	component: apiserver provider: kubernetes	10.96.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	33 minutes ago

Name	Namespace	Labels	Created
kube-root-ca.crt	default	-	33 minutes ago

Then click on Sign in. The dashboard is created.